

COSC 522 – Machine Learning

Bayes Decision Theory – In-depth Discussion (Parametric vs. Non-Parametric)

Hairong Qi, Gonzalez Family Professor
Electrical Engineering and Computer Science
University of Tennessee, Knoxville
<https://www.eecs.utk.edu/people/hairong-qi/>
Email: hqi@utk.edu

Bayes' Formula (Bayes' Rule)

conditional probability
density function (pdf)
or "likelihood"

from training data

prior probability (*a-*
priori probability)

from domain knowledge

$$P(w_j|x) = \frac{p(x|w_j) P(w_j)}{p(x)}$$

j index for different
classes
w_j: different classes
x: training sample

posterior probability
(*a-posteriori*
probability)

$$p(x) = \sum_{j=1}^c p(x|w_j) P(w_j)$$

normalization constant (evidence)

Part I

**In-Depth Discussion on Parametric
Learning: Discriminant Functions (Three
Cases with m-d Gaussian pdf)**

Bayes Decision Theory

$$P(\omega_j | x) = \frac{p(x | \omega_j) P(\omega_j)}{p(x)}$$

Maximum
Posterior
Probability

For a given x , if $P(\omega_1 | x) > P(\omega_2 | x)$,
then x belongs to class 1, otherwise, 2.

Discriminant
Function

The classifier will assign a feature vector x to class ω_i if
 $g_i(x) > g_j(x)$

Parametric
Learning
with
Gaussian
pdf

Case 1: Minimum Euclidean Distance (Linear Machine), $\Sigma_i = \sigma^2 I$

Case 2: Minimum Mahalanobis Distance (Linear Machine), $\Sigma_i = \Sigma$

Case 3: Quadratic classifier, $\Sigma_i = \text{arbitrary}$

All assuming Gaussian pdf

Multivariate Normal Density

$$p(\vec{x}) = \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp \left[-\frac{1}{2} (\vec{x} - \vec{\mu})^T \Sigma^{-1} (\vec{x} - \vec{\mu}) \right]$$

\vec{x} : d - component column vector

$\vec{\mu}$: d - component mean vector

Σ : d - by - d covariance matrix

$|\Sigma|$: determinant

Σ^{-1} : inverse

$$\vec{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_d \end{bmatrix}, \vec{\mu} = \begin{bmatrix} \mu_1 \\ \vdots \\ \mu_d \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} \sigma_{11} & \cdots & \sigma_{1d} \\ \vdots & \ddots & \vdots \\ \sigma_{d1} & \cdots & \sigma_{dd} \end{bmatrix} = \begin{bmatrix} \sigma_1^2 & \cdots & \sigma_{1d} \\ \vdots & \ddots & \vdots \\ \sigma_{d1} & \cdots & \sigma_d^2 \end{bmatrix}$$

When $d = 1$, $p(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp \left[-\frac{1}{2} \frac{(x - \mu)^2}{\sigma^2} \right]$

Estimating Normal Densities

◆ Calculate μ, Σ

$$\vec{\mu}_i = \begin{bmatrix} \mu_{i1} = \frac{1}{n_i} \sum_{k=1}^{n_i} x_{k1} \\ \vdots \\ \mu_{id} = \frac{1}{n_i} \sum_{k=1}^{n_i} x_{kd} \end{bmatrix}$$

$$\Sigma_i = \begin{bmatrix} \sigma_{11} & \cdots & \sigma_{1d} \\ \vdots & \ddots & \vdots \\ \sigma_{d1} & \cdots & \sigma_{dd} \end{bmatrix} = \frac{1}{n_i - 1} \sum_{k=1}^{n_i} (\vec{x}_k - \vec{\mu}_i)(\vec{x}_k - \vec{\mu}_i)^T$$

Covariance

For d sets of variates denoted $\{x_1\}, \dots, \{x_p\}, \dots, \{x_q\}, \dots, \{x_d\}$ the covariance $\sigma_{pq} = \text{cov}(x_p, x_q)$ of x_p and x_q is defined by

$$\begin{aligned}\text{cov}(x_p, x_q) &= E[(x_p - \mu_p)(x_q - \mu_q)] \\ &= E[x_p x_q] - E[x_p \mu_q] - E[\mu_p x_q] + E[\mu_p \mu_q] \\ &= E[x_p x_q] - \mu_q E[x_p] - \mu_p E[x_q] + \mu_p \mu_q \\ &= E[x_p x_q] - \mu_q \mu_p - \mu_p \mu_q + \mu_p \mu_q \\ &= E[x_p x_q] - \mu_q \mu_p\end{aligned}$$

$$\begin{aligned}\text{When } p = q, \sigma_{pp} &= \text{cov}(x_p, x_p) = E[x_p x_p] - \mu_p \mu_p \\ &= E[x_p^2] - (E[x_p])^2 \\ &= \sigma_p^2\end{aligned}$$

Discriminant Function

$$P(\omega_j|x) = \frac{p(x|\omega_j) P(\omega_j)}{p(x)}$$

- ◆ One way to represent pattern classifier - use discriminant functions $g_i(x)$

$$g_i(x) = P(\omega_i|x)$$

$$g_i(x) = p(x|\omega_i)P(\omega_i)$$

$$g_i(x) = \ln p(x|\omega_i) + \ln P(\omega_i)$$

The classifier will assign a feature vector x to class ω_i if

$$g_i(x) > g_j(x)$$

- ◆ For two-class cases,

$$g(x) = g_1(x) - g_2(x) = P(\omega_1 | x) - P(\omega_2 | x)$$

Discriminant Function for Normal Density

$$p(\vec{x} | w) = \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp \left[-\frac{1}{2} (\vec{x} - \vec{\mu})^T \Sigma^{-1} (\vec{x} - \vec{\mu}) \right]$$

$$\begin{aligned} g_i(\vec{x}) &= \ln p(\vec{x} | \omega_i) + \ln P(\omega_i) \\ &= -\frac{1}{2} (\vec{x} - \vec{\mu}_i)^T \Sigma_i^{-1} (\vec{x} - \vec{\mu}_i) - \frac{d}{2} \ln(2\pi) - \frac{1}{2} \ln |\Sigma_i| + \ln P(\omega_i) \\ &= -\frac{1}{2} (\vec{x} - \vec{\mu}_i)^T \Sigma_i^{-1} (\vec{x} - \vec{\mu}_i) - \frac{1}{2} \ln |\Sigma_i| + \ln P(\omega_i) \end{aligned}$$

Case 1: $\Sigma_i = \sigma^2 I$

- ◆ The features are statistically independent, and have the same variance
- ◆ Geometrically, the samples fall in equal-size hyperspherical clusters
- ◆ Decision boundary: hyperplane of $d-1$ dimension

$$\Sigma = \begin{bmatrix} \sigma^2 & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \sigma^2 \end{bmatrix}, |\Sigma| = \sigma^{2d}, \Sigma^{-1} = \begin{bmatrix} \frac{1}{\sigma^2} & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \frac{1}{\sigma^2} \end{bmatrix}$$

Linear Discriminant Function and Linear Machine

$\|\vec{x} - \vec{\mu}_i\|$: the Euclidean norm (distance)

$$\|\vec{x} - \vec{\mu}_i\|^2 = (\vec{x} - \vec{\mu}_i)^T (\vec{x} - \vec{\mu}_i)$$

$$\begin{aligned} g_i(\vec{x}) &= -\frac{\|\vec{x} - \vec{\mu}_i\|^2}{2\sigma^2} + \ln P(\omega_i) \\ &= -\frac{\vec{x}^T \vec{x} - 2\vec{\mu}_i^T \vec{x} + \vec{\mu}_i^T \vec{\mu}_i}{2\sigma^2} + \ln P(\omega_i) \end{aligned}$$

$$g_i(\vec{x}) = \frac{\vec{\mu}_i^T}{\sigma^2} \vec{x} - \frac{\vec{\mu}_i^T \vec{\mu}_i}{2\sigma^2} + \ln P(\omega_i)$$

Minimum-Distance Classifier

- ◆ When $P(\omega_i)$ are the same for all c classes, the discriminant function is actually measuring the minimum distance from each x to each of the c mean vectors

$$g_i(\vec{x}) = -\frac{\|\vec{x} - \vec{\mu}_i\|^2}{2\sigma^2}$$

Case 2: $\Sigma_i = \Sigma$

- ◆ The covariance matrices for all the classes are identical but not a scalar of identity matrix.
- ◆ Geometrically, the samples fall in hyperellipsoidal
- ◆ Decision boundary: hyperplane of d-1 dimension

$$\begin{aligned} g_i(\vec{x}) &= \ln p(\vec{x} | \omega_i) + \ln P(\omega_i) \\ &= -\frac{1}{2} \left(\vec{x} - \vec{\mu}_i \right)^T \Sigma_i^{-1} \left(\vec{x} - \vec{\mu}_i \right) + \ln P(\omega_i) \\ &= \vec{\mu}_i^T \left(\Sigma^{-1} \right)^T \vec{x} - \frac{1}{2} \vec{\mu}_i^T \Sigma^{-1} \vec{\mu}_i + \ln P(\omega_i) \end{aligned}$$

Squared Mahalanobis distance

Case 3: $\Sigma_i = \text{arbitrary}$

- ◆ The covariance matrices are different from each category
- ◆ Quadratic classifier
- ◆ Decision boundary: hyperquadratic for 2-D Gaussian

$$\begin{aligned} g_i(\vec{x}) &= \ln p(\vec{x} | \omega_i) + \ln P(\omega_i) \\ &= -\frac{1}{2} (\vec{x} - \vec{\mu}_i)^T \Sigma_i^{-1} (\vec{x} - \vec{\mu}_i) - \frac{1}{2} \ln |\Sigma_i| + \ln P(\omega_i) \\ &= -\frac{1}{2} \vec{x}^T \Sigma_i^{-1} \vec{x} + \vec{\mu}_i^T (\Sigma_i^{-1}) \vec{x} - \frac{1}{2} \vec{\mu}_i^T \Sigma_i^{-1} \vec{\mu}_i - \frac{1}{2} \ln |\Sigma_i| + \ln P(\omega_i) \end{aligned}$$

Questions

- What is a discriminant function?
- What is a multivariate Gaussian (or normal density function)?
- What is the covariance matrix and what is its dimension?
- What would the covariance matrix look like if the features are independent from each other?
- What would the covariance matrix look like if the features are independent from each other AND have the same spread in each dimension?
- What is minimum (Euclidean) distance classifier? Is it a linear or quadratic classifier (machine)? What does the decision boundary look like?
- What are the assumptions made when using a minimum (Euclidean) distance classifier?
- What is minimum (Mahalanobis) distance classifier? Is it a linear or quadratic classifier (machine)? What does the decision boundary look like?
- What are the assumptions made when using a minimum (Mahalanobis) distance classifier?
- What does the decision boundary look like for a quadratic classifier?
- What are the cost functions for the discriminant functions? And what is the optimization method used to find the best solution?

Part II

In-Depth Discussion on Non-Parametric Learning: Why kNN?

kNN in Classification

$$p_n(x) = \frac{k_n/n}{V}$$

- Given c training sets from c classes, the total number of samples is

$$n = \sum_{m=1}^c n_m$$

- Given a point \mathbf{x} at which we wish to determine the statistics, we find the hypersphere of volume V which just encloses k points from the combined set. If within that volume, k_m of those points belong to class m , then we estimate the density for class m by

$$p(x|\omega_m) = \frac{k_m/n_m}{V} \quad P(\omega_m) = \frac{n_m}{n} \quad p(x) = \frac{k/n}{V}$$

$$P(\omega_m | x) = \frac{p(x|\omega_m)P(\omega_m)}{p(x)} = \frac{\frac{k_m}{n_m} \frac{n_m}{n}}{\frac{k}{nV}} = \frac{k_m}{k}$$

kNN Decision Boundary




Figure 2.28 Plot of 200 data points from the oil data set showing values of x_6 plotted against x_7 , where the red, green, and blue points correspond to the ‘laminar’, ‘annular’, and ‘homogeneous’ classes, respectively. Also shown are the classifications of the input space given by the K -nearest-neighbour algorithm for various values of K .

From [Bishop 2006]

Parzen Windows

$$p_n(x) = \frac{k_n/n}{V}$$

- The density estimation at x is calculated by counting the number of samples fall within a hypercube of volume V centered at x
- Let R be a d -dimensional hypercube, whose edges are h units long. Its volume is then $V=h^d$
- Introducing the “window” function

$$\varphi(u) = \begin{cases} 1 & |u_j| \leq 0.5 \quad j=1, \dots, d \\ 0 & \text{otherwise} \end{cases}$$

- Calculate k_n


$$k_n = \sum_{i=1}^n \varphi\left(\frac{x - x_i}{h}\right)$$

- Hence

$$p_n(x) = \frac{1}{n} \sum_{i=1}^n \frac{\varphi\left(\frac{x - x_i}{h}\right)}{V}$$

Problems of Parzen Windows

- Hypercube – why should a point just inside the hypercube contribute the same as a point very near to \mathbf{x} , while a point just outside the hypercube contributes nothing? – Use a **continuous** window function



- How to choose h ? – Depend on **the number of samples**.
- ... but the hypercube is of **fixed** volume!
- How does kNN solve the “fixed” volume problem?

$$h = \frac{1}{\sqrt{n}}$$

Potential Issues of kNN

- What is a good value of “k”? $k_n = \sqrt{n}$
- What kind of distance should be used to measure “nearest”
 - Euclidean metric is a reasonable measurement
- Computation burden
 - Massive storage burden
 - Need to compute the distance from the unknown to all the neighbors