



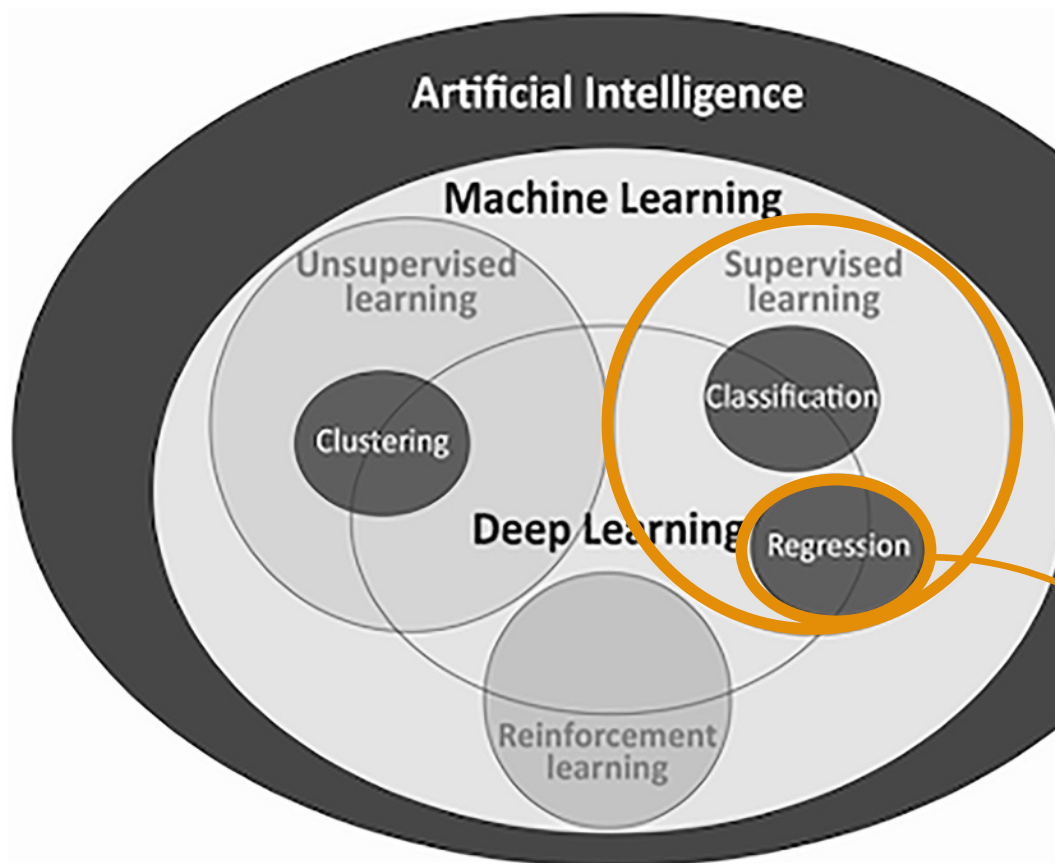
THE UNIVERSITY OF
TENNESSEE
KNOXVILLE

**AICIP
RESEARCH**

COSC 522 – Machine Learning

Regression

Hairong Qi, Gonzalez Family Professor
Electrical Engineering and Computer Science
University of Tennessee, Knoxville
<https://www.eecs.utk.edu/people/hairong-qi/>
Email: hqi@utk.edu



Date & Lecture Recordings	Topics
Part 1: Statistical Methods	
Bayesian Learning	
08/20 (T)	Introduction
08/22 (R)	Bayesian Decision Theory and Parametric Learning
08/27 (T)	Bayesian Decision Theory and Non-Parametric Learning
08/29 (R)	Case Study: Representation for Natural Language (taught by An Cozma)
09/03 (T)	Parametric vs. Non-Parametric Learning: Some In-Depth Discus
09/05 (R)	Homework and Project Discussion (taught by Fanqi Wang)
Neural Networks	
09/10 (T)	Biological Neuron and Perceptron
09/12 (R)	Perceptron
09/17 (T)	Back Propagation and Gradient Descent
09/19 (R)	Back Propagation
09/20 (F)	TRUST-AI Seminar
09/24 (T)	Kernel Methods and Review
09/26 (R)	Test 1
10/01 (T)	Kernel Methods and Support Vector Machine
Regression	
10/03 (R)	Logistic Regression
10/08 (T)	Fall Break (No Class)

Questions

- Classification vs. Regression vs. Generation
- Linear regression and various basis functions
- What is global vs. local basis function?
- Maximum likelihood and least-square solution
- Least-square with regularization

PART I: REGRESSION AND BASIS FUNCTIONS

Linear regression (Linear function in \mathbf{w})

Generally

$$y(\mathbf{x}, \mathbf{w}) = \sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x})$$

where $\phi_j(\mathbf{x})$ are known as *basis functions*.

Typically, $\phi_0(\mathbf{x}) = 1$, so that w_0 acts as a bias.

In the simplest case, we use linear basis functions :


$$\phi_d(\mathbf{x}) = x_d.$$

Basis function - Polynomial

Polynomial basis functions:

$$\phi_j(x) = x^j.$$

These are **global**; a small change in x affect all basis functions.




Polynomial regression

Basis function - Gaussian

Gaussian basis functions:

$$\phi_j(x) = \exp \left\{ -\frac{(x - \mu_j)^2}{2s^2} \right\}$$

These are **local**; a small change in x only affect nearby basis functions. μ and s control location and scale (width).



Basis function - Sigmoid


Sigmoidal basis functions:

$$\phi_j(x) = \sigma\left(\frac{x - \mu_j}{s}\right)$$

where

$$\sigma(a) = \frac{1}{1 + \exp(-a)}.$$

Also these are **local**; a small change in x only affect nearby basis functions. μ and s control location and scale (slope).



Logistic regression

- The logistic function

$$\phi(x_k) = p(x_k) = \sigma\left(\frac{x_k - \mu}{s}\right) = \frac{1}{1 + e^{-\frac{x_k - \mu}{s}}}$$

- The log loss for the k th point

$$\begin{cases} -\ln \phi_k & \text{if } y_k = 1 \\ -\ln (1 - \phi_k) & \text{if } y_k = 0 \end{cases}$$

- The cost function: cross entropy

$$l(\beta_0, \beta_1) = \sum_k -y_k \log p_k - (1 - y_k) \log(1 - p_k)$$

- Find μ and s that best predict the probability of x belonging to a certain category

PART II: HOW TO SOLVE?

Maximum Likelihood and Least Squares (1)

Assume observations from a deterministic function with added Gaussian noise:

$$t = y(\mathbf{x}, \mathbf{w}) + \epsilon \quad \text{where} \quad p(\epsilon|\beta) = \mathcal{N}(\epsilon|0, \beta^{-1})$$

which is the same as saying,

$$p(t|\mathbf{x}, \mathbf{w}, \beta) = \mathcal{N}(t|y(\mathbf{x}, \mathbf{w}), \beta^{-1}).$$

Given observed inputs, $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$, and targets, $\mathbf{t} = [t_1, \dots, t_N]^T$, we obtain the likelihood function

$$p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n|\mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1}).$$

Maximum Likelihood and Least Squares (2)

Taking the logarithm, we get

$$\begin{aligned}\ln p(\mathbf{t}|\mathbf{w}, \beta) &= \sum_{n=1}^N \ln \mathcal{N}(t_n | \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n), \beta^{-1}) \\ &= \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi) - \beta E_D(\mathbf{w})\end{aligned}$$

where

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n)\}^2$$

is the sum-of-squares error.

Maximum Likelihood and Least Squares (3)

Computing the gradient and setting it to zero yields

$$\nabla_{\mathbf{w}} \ln p(\mathbf{t}|\mathbf{w}, \beta) = \beta \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\} \phi(\mathbf{x}_n)^T = \mathbf{0}.$$

Solving for \mathbf{w} , we get

$$\mathbf{w}_{\text{ML}} = \left(\Phi^T \Phi \right)^{-1} \Phi^T \mathbf{t}$$

The Moore-Penrose pseudo-inverse, Φ^\dagger .

where

$$\Phi = \begin{pmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \cdots & \phi_{M-1}(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & \phi_1(\mathbf{x}_2) & \cdots & \phi_{M-1}(\mathbf{x}_2) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}_N) & \phi_1(\mathbf{x}_N) & \cdots & \phi_{M-1}(\mathbf{x}_N) \end{pmatrix}.$$

Least squares with regularization terms

- L2 norm (sum of square or Weight decay): Ridge regression
- L1 norm (LASSO regression), $q=1$ (sparsity)
- L12 norm (ElasticNet regression), $q=1$ and 2 , $M=2$.

$$\frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}.$$

$$\mathbf{w} = (\lambda \mathbf{I} + \Phi^T \Phi)^{-1} \Phi^T \mathbf{t}.$$

$$\frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(\mathbf{x}_n)\}^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|^q$$





Figure 3.3 Contours of the regularization term in (3.29) for various values of the parameter q .

Figure 3.4 Plot of the contours of the unregularized error function (blue) along with the constraint region (3.30) for the quadratic regularizer $q = 2$ on the left and the lasso regularizer $q = 1$ on the right, in which the optimum value for the parameter vector \mathbf{w} is denoted by \mathbf{w}^* . The lasso gives a sparse solution in which $w_1^* = 0$.



Linear regression - Summary

- Linear regression
 - Simple linear regression (d=1)
 - Multiple linear regression (d>1)
 - Polynomial regression $\phi_j(x) = x^j$.
 - Logistic regression $\phi_j(x) = \sigma\left(\frac{x - \mu_j}{s}\right)$
- Solving linear regression with maximum likelihood
 - Unconstrained formulation leads to least-squares solution

$$\mathbf{w}_{\text{ML}} = \left(\Phi^T \Phi\right)^{-1} \Phi^T \mathbf{t}$$
 - Constrained formulation with regularization terms
 - L2 norm \rightarrow Ridge regression (q=2)
 - L1 norm \rightarrow LASSO regression (q=1)
 - L12 norm \rightarrow ElasticNet regression (q=1 and q=2)